

Discovery

Michael McCool

2020 June 22

Outline

- Requirements
- Key Design Decisions
- Two-Phase Architecture
 - Privacy Considerations
 - Introduction
 - Exploration
- Key Issues for Directory Service
- Other Open Issues

Discovery Requirements

- Capabilities
 - Support both local and global/remote discovery (unconstrained by network domain)
 - Support “localizable” discovery (constrainable by location)
 - Support some form of “semantic query”
 - Be usable by constrained devices
 - Support both
 - Directory services for searching large repositories of Things
 - Peer-to-peer (self-identifying) discovery
- Privacy-Preserving Architecture
 - Respect device and information Lifecycle
 - Distribute TDs only to authenticated and authorized users
 - Don’t leak private data to unauthorized users
 - Don’t leak information usable to INFER private information to unauthorized users
- Alignment with existing standards
 - E.g. IETF CoRE Resource Directories, CoRE Link Format, DNS-SD, DID, ...
 - Align with WoT Scripting API

Key Design Decisions

See: <https://github.com/w3c/wot-discovery/blob/master/design.md>

Resolution: Use Two-Phase Architecture

- Introduction, Exploration

Under Discussion:

- Support as Introductions:
 - DHCP, DNS-SD, DID Documents, QR Codes, EddyStone Beacons.
 - Probably also: RFID, CoRE RD
 - Need consistent set of "link types" corresponding to each exploration type
- Support as Explorations:
 - Well-Known Location on Device
 - Standardized Directory Service

Two-Phase Architecture

1. Introduction

- “First Contact” Protocol
 - Answers the question: how to start?
- Open
 - Can be accessed with no or limited access controls
- Lightweight
 - Does not use significant resources on responder
 - Resistant to Denial of Service attacks (limited/finite resources to support)
- Provides intentionally limited information
 - Avoid leaking any metadata that can be used to infer private data
 - This includes types of devices, device ids, owners, timestamps, etc.

2. Exploration

- Authentication and authorization required
- Supports more complex search capabilities
- Provides access to rich metadata
- Access controls can limit data returned

Privacy Issues

- Two-phase approach ***not sufficient*** to preserve privacy in all contexts
 - Privacy preservation also depends on the design of API
 - API needs to hide data that can be used to infer private information, such as location of device doing the discovery
- Third-party code context (eg browser):
 - If discovery API follows two-phase structure, where Introduction returns list of directories, then even without authenticating the list of directories visible can possibly be used to infer location via fingerprinting
 - This is especially true if the discovery mechanism can be constrained to particular Introduction mechanisms.
 - May also be a problem in proposed non-browser contexts, eg. “Edge Workers”

Introduction

- “First Contact” protocol
 - Output: Address of exploration service, for example, a directory service
 - Need not be broadcast; could use well-known network services (eg DNS, DHCP, etc)
- Address should not leak any other metadata, e.g. type of devices
- Can have multiple mechanisms for introduction
 - Local: QR code, mDNS, DNS-SD, DHCP, Bluetooth beacons (Eddystone), etc.
 - Global: Search engine (incl. spatial queries), well-known global repositories, company repositories, cities, etc.
 - Self: Well-known addresses, eg “.well-known/wot-td”
- Existing mechanisms that have lists of typed links can also be used here:
 - CoRE RD, DID Documents, DNS-SD, etc.
 - Use these to find directories rather than to distribute metadata directly
- *May* in some cases point directly at a Thing Description
 - Degenerate case: Well-known address like a “directory” that has only one TD
 - Still requires authentication in principle to access content

Exploration

- Authentication required...
 - Then provides access to one or more TDs
- Simple: Self-hosted self-identification
 - Single TD retrievable by an authenticated GET
 - From direct target of URL
 - Perhaps also in well-known location, e.g. {base}/.well-known/wot-td
- Advanced: Queryable Directory service
 - Lightweight: specific query parameters, eg. location, keywords
 - Middle: JSONPath and/or XPath query language
 - Full: (sub-)SPARQL semantic query, GraphQL, etc.


Key Issues for Directory Service

- [What form of Query?](#)
 - May be multiple levels
 - SPARQL makes sense for the high end (semantic query), but an intermediate level (syntactic query) would be useful for local gateways, etc.
- Other [Design Issues:](#)
 - Directory service as a Thing with a TD describing its API:
 - Self-describing using well-known location approach
 - Directories can then index other directories; how to summarize?
 - TCP/HTTP baseline protocol
 - Support for Partial TDs and/or TD Fragments (has implications for JSON-LD)
 - Pagination (when returning multiple TDs)
 - Out-of-band metadata
 - TD modification and update notifications

Other Open Issues

- [Support for Signed TDs](#)
 - Signed TDs prevent modification of the TD
 - Directories need way to [support out-of-band information](#)
- [Returning Partial TDs in Query Results](#)
 - Potentially breaks JSON-LD contexts
- [Handling TTL](#)
 - What limits should be placed on TTL? Need to consider use cases.
- [Mutable TDs and Notifications](#)
 - Should directories notify subscribers if a TD changes?
- [Profiles for Directories](#)
 - Limiting the maximum length of a TD

Backup



Resources

- Repository: <https://github.com/w3c/wot-discovery>
- Prior work: <https://github.com/w3c/wot-discovery/blob/master/prior.md>
- Background: <https://github.com/w3c/wot-discovery/blob/master/background.md>
- Requirements: <https://github.com/w3c/wot-discovery/blob/master/requirements.md>
- Design: <https://github.com/w3c/wot-discovery/blob/master/design.md>
- Proposals: <https://github.com/w3c/wot-discovery/tree/master/proposals>

Current Scripting API – Discussion Points

- Allows discovery mechanism to be specified
- Returns discovered TDs
- Assumes authentication/authorization is handled out-of-band

Discussion Points:

- Not incompatible with two-phase approach
 - Authentication/authorization needs to be set up outside script
- Query format needs to be better standardized
 - What query forms are supported? What are the parameters?
 - How can I search for Things in a particular physical location (that I am not necessarily at?)
- Options to select mechanisms may be a privacy risk
 - If I can discover a thing via Bluetooth, I know it is within a few meters
 - Perhaps the mechanisms and desired location should also be specified out of band
- Alternatives:
 - “Discovery sub-API” could be factored out and run during setup, like security config
 - Things to be consumed by a script perhaps managed declaratively using dependences